

fun-ai-talks

Let's read 7 DeepMind Alpha[*] papers!

AlphaGo

2016

AlphaGo Zero

2017

AlphaZero

2018

AlphaStar

2019

AlphaFold 2

2021

AlphaCode

2022

AlphaTensor

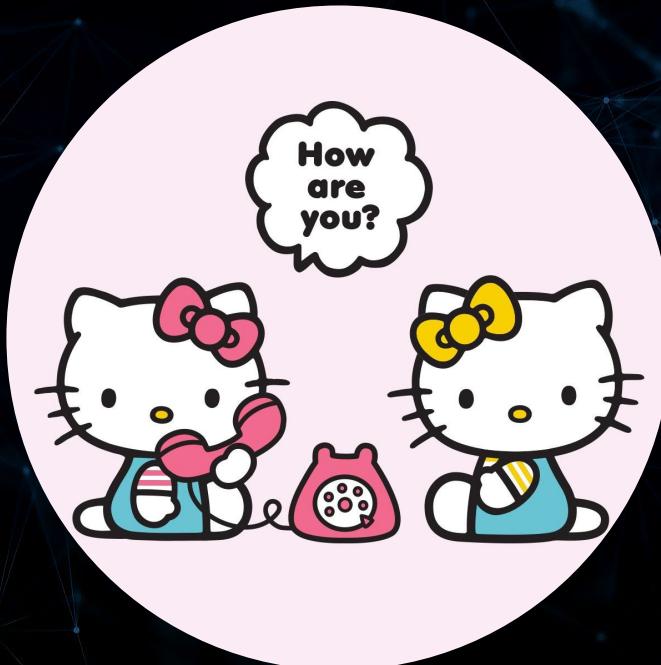
hululu.zhu@gmail.com

10/2022

Disclaimer

- All content in this deck is based on publicly published research papers
- All opinions in this slide deck are my personal own, and not those of DeepMind or Google

How are you today?



About me, 10+ years at Google

My ML work experience starts at
2016 at Google Fiber



intern

SWE

Why I create this course?

- DeepMind published breakthrough research since 2015
 - Some of the best work is named after *Alpha** pattern
 - It seems quite meaningful to combine into one shared talk!
 - To my best knowledge, nobody has made a course/deck to summarize all the *Alpha** papers as of 05/2022, so I will try
 - Updated 10/2022 to add recent AlphaTensor
- And sharing is fun!
 - Sharing and simplification for more in-depth learning
 - [go/fun-ai-courses](#) is my personal project for ML education

Quick glance of the 7 Alpha[*] papers

- **AI for game**
 - 2016 AlphaGo: AI for Go with human data
 - 2017 AlphaGo Zero: AI for Go without human data
 - 2018 AlphaZero: A general AI framework for Go, Chess, Shogi
 - 2019 AlphaStar: Multi-agent AI for Starcraft
- **AI for science/engineering**
 - 2021 AlphaFold 2: AI to predict protein folding 3d structure
 - 2022 AlphaCode: AI to solve competitive coding problems
 - 2022 AlphaTensor: AI to improve matrix multiplication algorithms

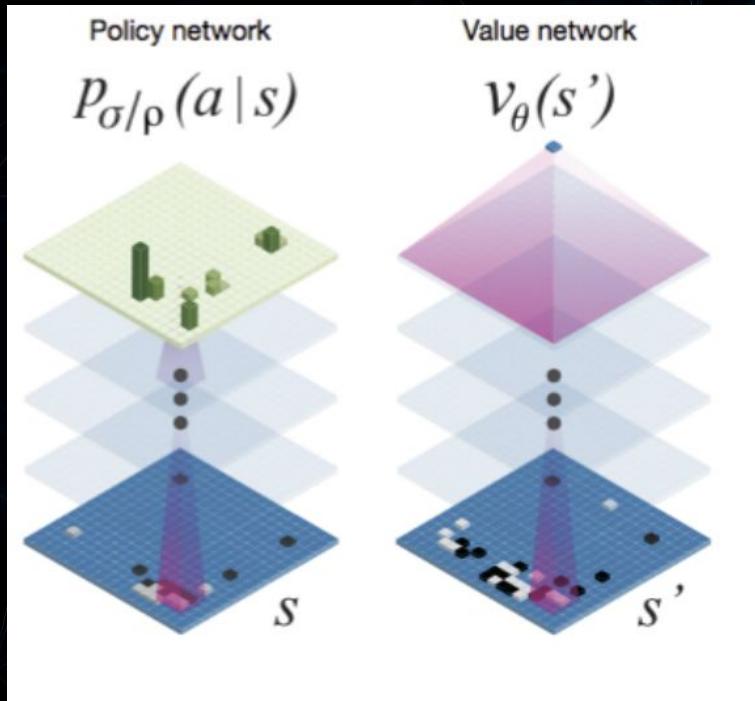
* The years are the publication years, actual research time may start much earlier than that

AlphaGo 2016



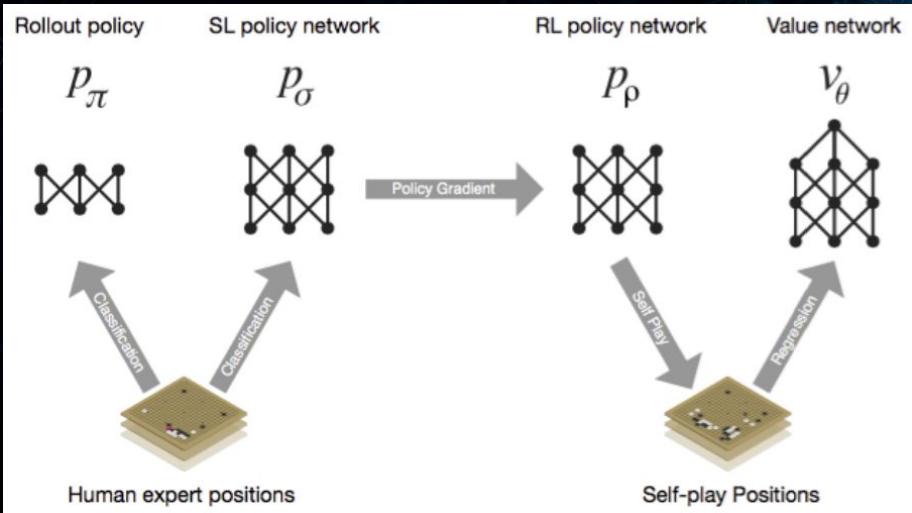
- AI for Go game is hard!
- 2 types of networks
 - Value Network
 - Policy Network
 - Monte Carlo Tree Search to effectively combine value and policy networks
- 2 learnings applied
 - supervised learning
 - reinforcement learning

AlphaGo 2016, 2 networks: policy and value



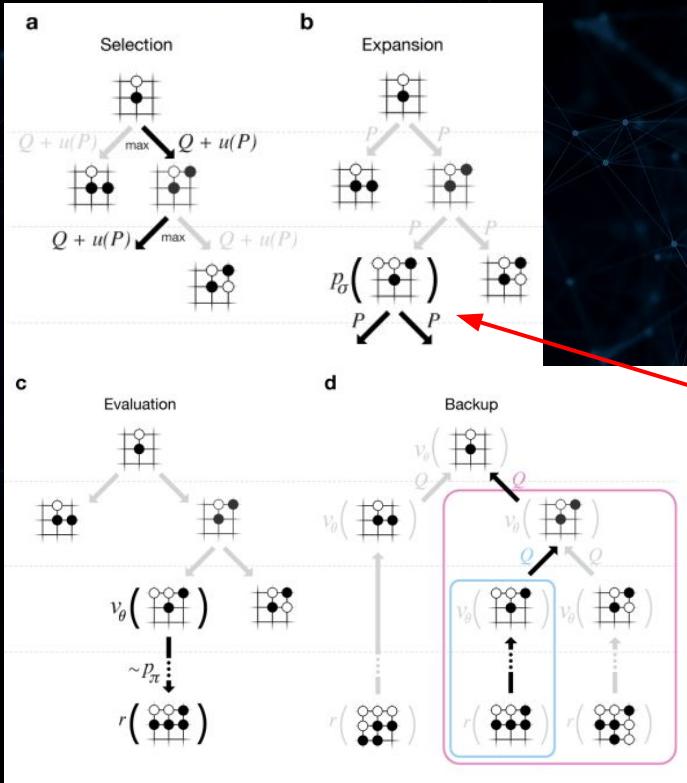
- CNN Policy Network
 - Input: Encoded go board state features s (more than positions)
 - Output: Prob distribution of action a
- CNN Value Network
 - Input: Encoded go board state features s
 - Output: A single scalar value to predict win or loss

AlphaGo 2016, Supervised Learning (SL) & Reinforcement learning (RL)



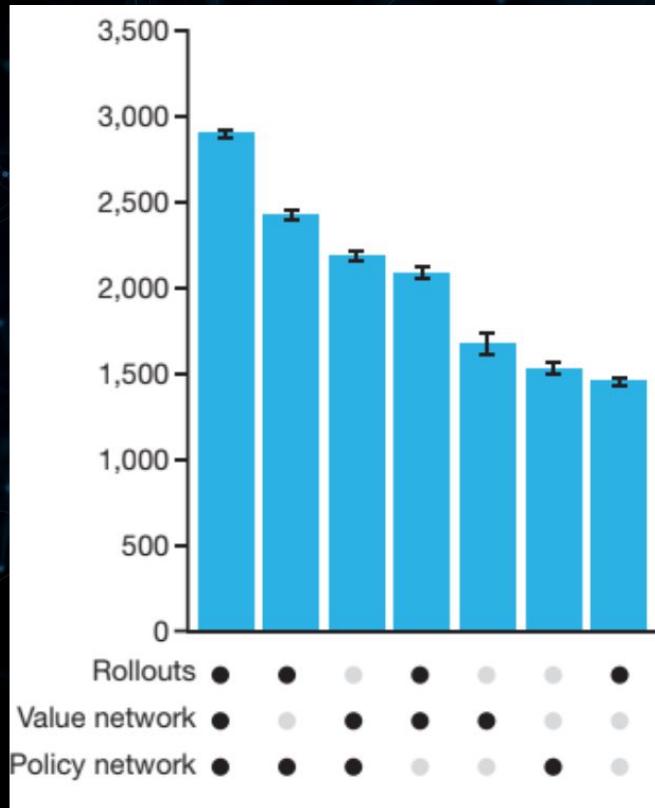
- Fast Rollout Policy network
 - A “simpler, less accurate” but 1000x faster network to imitate expert actions
- SL policy network
 - An accurate but slower network to better imitate expert actions
- RL policy network (self-play RL)
 - Initiated from SL policy network, optimize for “win” by self-play!
- Value network
 - Predict outcome from the RL policy network self-play

AlphaGo 2016, Monte Carlo Tree Search to “lookahead search”

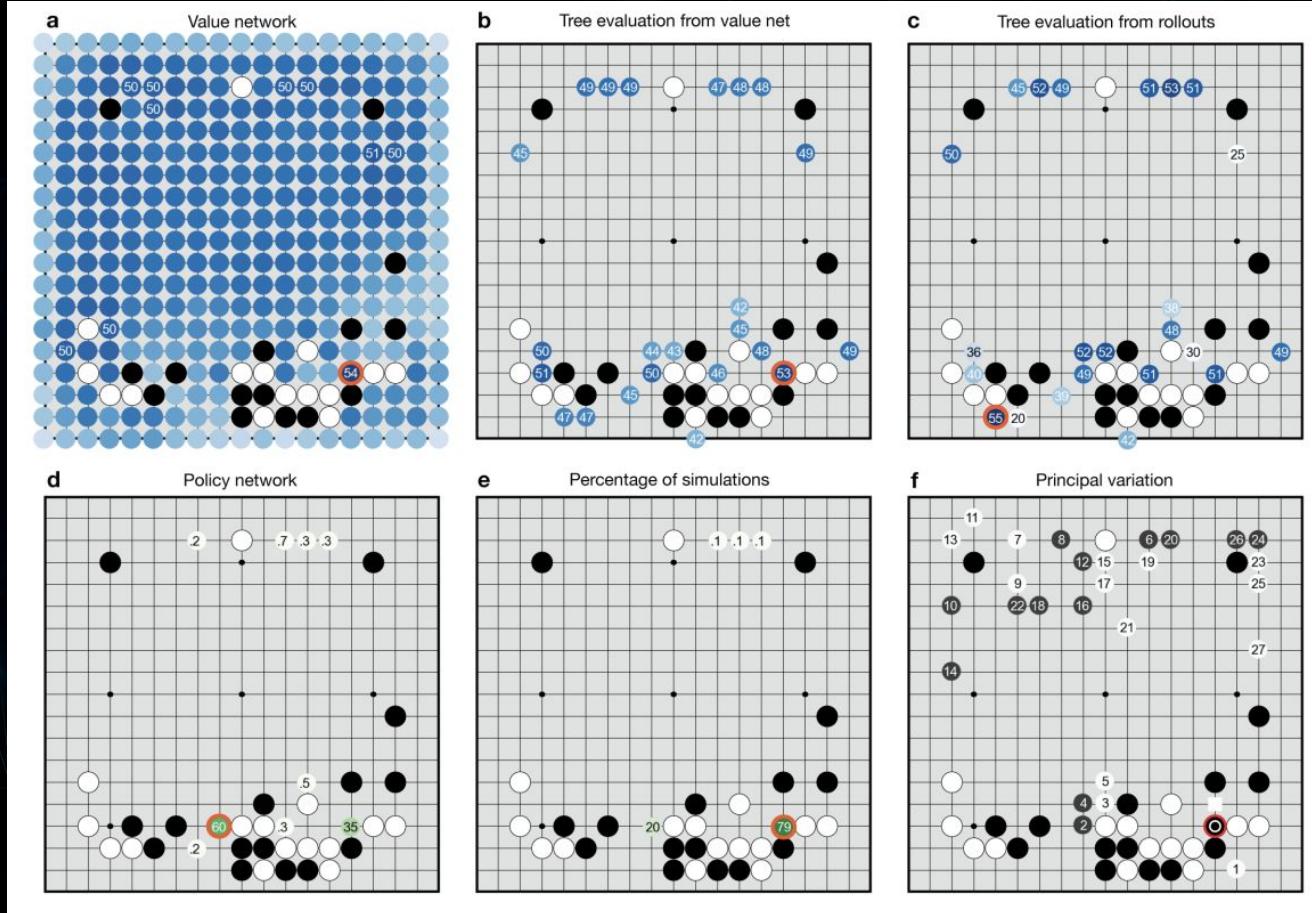


- MCTS is efficient to sample and approximate “optimal policy”
- Combined value network (predict win/loss) and fast rollout outcome (fast simulate to report win/loss)
 - Discounted by repeated visits to encourage “exploration”
- Note the supervised learning policy network is used to get prior action distribution!
 - SL network to imitate experts may “select a diverse beam of promising moves”, while RL network “optimized for single best move”

AlphaGo 2016, ablation analysis of different networks, all 3 networks are necessary!



AlphaGo 2016, example how MCTS works



AlphaGo Zero 2017

The screenshot shows the 'nature' journal website. At the top, there is a search bar and a login link. Below the header, there are navigation links for 'Explore content', 'About the journal', and 'Publish with us'. The main content area displays the following information:

Published: 19 October 2017

Mastering the game of Go without human knowledge

David Silver Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel & Demis Hassabis

Nature 550, 354–359 (2017) | [Cite this article](#)

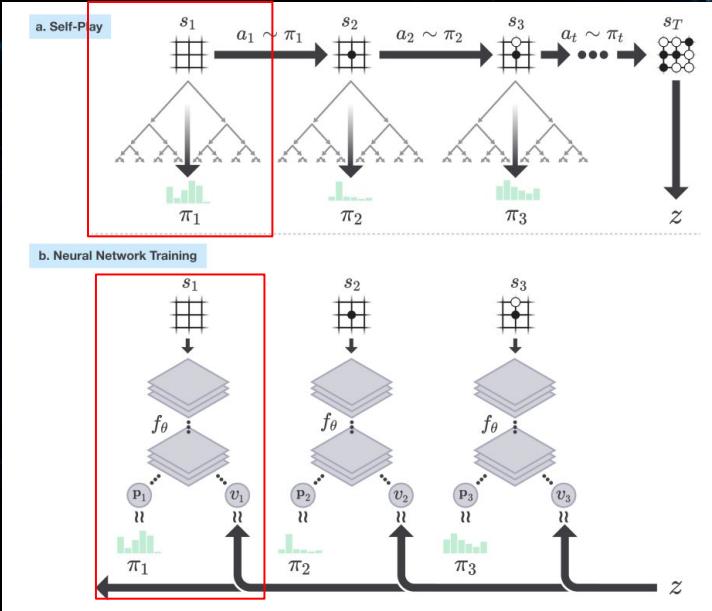
319k Accesses | 3103 Citations | 2570 Altmetric | [Metrics](#)

Abstract

A long-standing goal of artificial intelligence is an algorithm that learns, *tabula rasa*, superhuman proficiency in challenging domains. Recently, AlphaGo became the first program to defeat a world champion in the game of Go. The tree search in AlphaGo evaluated positions and selected moves using deep neural networks. These neural networks were trained by supervised learning from human expert moves, and by reinforcement learning from self-play. Here we introduce an algorithm based solely on reinforcement learning, without human data, guidance or domain knowledge beyond game rules. AlphaGo becomes its own teacher: a neural network is trained to predict AlphaGo's own move selections and also the

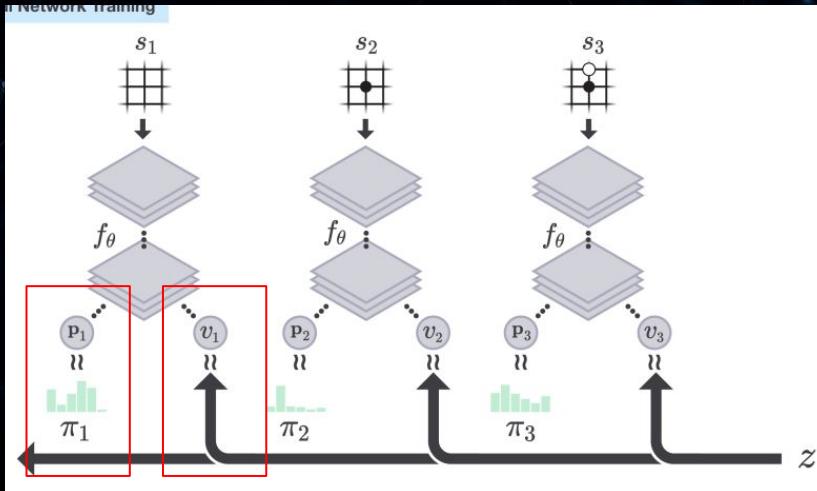
- Simpler than AlphaGo
 - No human expert data!
 - Only board info as input
 - Shared value/policy network to predict win/loss and action distribution
- Monte Carlo Tree Search merged into training
 - AlphaGo used MCTS with model inference only

AlphaGo Zero 2017, network and MCTS



- CNN + Residual block (resnet)
- Shared network outputs v_t (win prob), and ρ_t (action distribution)
- MCTS applied at t during self-play
 - Action dist ρ_t to expand tree
 - Value output v_t to predict win/loss
 - Combined scores by discounting visits to encourage exploration
 - Sampling is applied on MCTS in self-play (to decide a_t) instead of choosing max value
- Game stops at s_t with outcome z

AlphaGo Zero 2017, loss and training



$$\text{loss} = (z - v)^2 - \pi^\top \log \rho + c \|\theta\|^2$$

- “More accurately predict win/loss”
 - minimize $(z - v)^2$ so model can predict outcome more accurately
- “More align with the MCTS self-play algorithm”
 - Cross entropy $(-\pi^\top \log \rho)$ so that ρ_t approximate MCTS distribution π_t , with a temperature param
- “Smaller weights”, weight decay for $(c \|\theta\|^2)$

AlphaGo Zero 2017, training with “selected self play games” with “best player”!

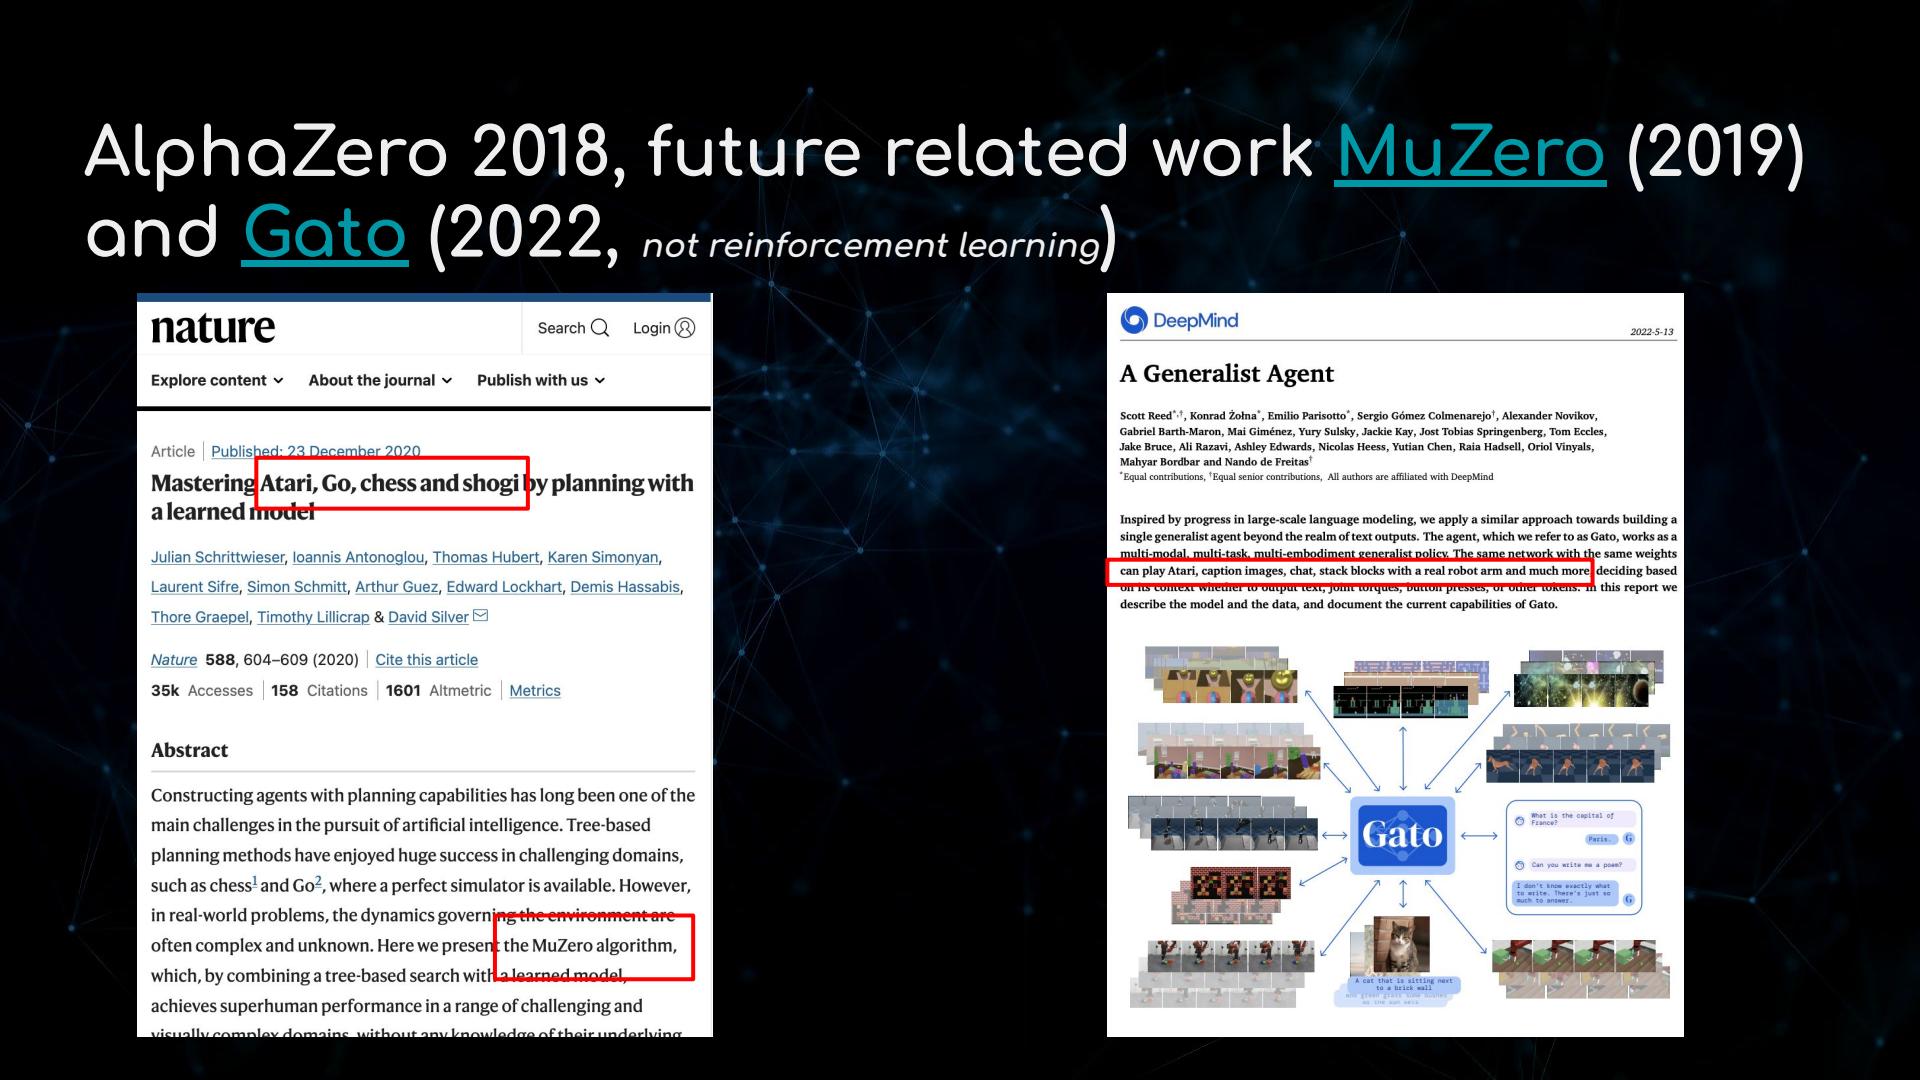
Evaluator To ensure we always generate the best quality data, we evaluate each new neural network checkpoint against the current best network f_{θ_*} before using it for data generation. The neural network f_{θ_i} is evaluated by the performance of an MCTS search α_{θ_i} that uses f_{θ_i} to evaluate leaf positions and prior probabilities (see Search Algorithm). Each evaluation consists of 400 games, using an MCTS with 1,600 simulations to select each move, using an infinitesimal temperature $\tau \rightarrow 0$ (i.e. we deterministically select the move with maximum visit count, to give the strongest possible play). If the new player wins by a margin of $> 55\%$ (to avoid selecting on noise alone) then it becomes the best player α_{θ_*} , and is subsequently used for self-play generation, and also becomes the baseline for subsequent comparisons.

AlphaZero 2018



- The same AI architecture can play chess, shogi, & Go
- Same arch as AlphaGo Zero
 - Same shared network to predict value and policy
- Different from AlphaGo Zero
 - Different size of board
 - No rotation/reflection
 - Non-binary outcome
 - “Continuous update” during self play, vs AlphaGo Zero only optimized on “best model” that won 55%+ over previous versions

AlphaZero 2018, future related work MuZero (2019) and Gato (2022, *not reinforcement learning*)



nature Search Login

Explore content ▾ About the journal ▾ Publish with us ▾

Article | Published: 23 December 2020

Mastering Atari, Go, chess and shogi by planning with a learned model

Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap & David Silver 

Nature 588, 604–609 (2020) | [Cite this article](#)

35k Accesses | 158 Citations | 1601 Altmetric | [Metrics](#)

Abstract

Constructing agents with planning capabilities has long been one of the main challenges in the pursuit of artificial intelligence. Tree-based planning methods have enjoyed huge success in challenging domains, such as chess¹ and Go², where a perfect simulator is available. However, in real-world problems, the dynamics governing the environment are often complex and unknown. Here we present the MuZero algorithm, which, by combining a tree-based search with a learned model, achieves superhuman performance in a range of challenging and visually complex domains, without any knowledge of their underlying

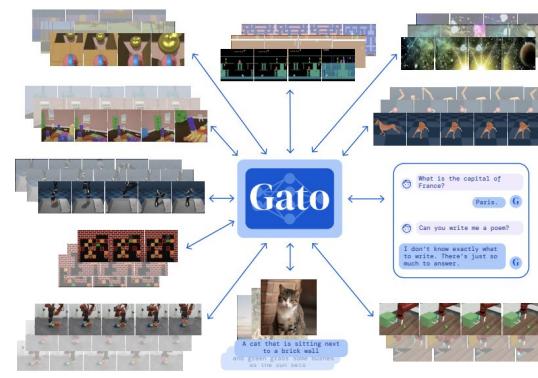
DeepMind 2022-5-13

A Generalist Agent

Scott Reed^{1,2}, Konrad Żolna³, Emilio Parisotto⁴, Sergio Gómez Colmenarejo¹, Alexander Novikov, Gabriel Barth-Maron, Mai Giménez, Yury Sulsky, Jackie Kay, Jost Tobias Springenberg, Tom Eccles, Jake Bruce, Ali Razavi, Ashley Edwards, Nicolas Heess, Yutian Chen, Raia Hadsell, Oriol Vinyals, Mahyar Bordbar and Nando de Freitas⁵

¹Equal contributions, ²Equal senior contributions, All authors are affiliated with DeepMind

Inspired by progress in large-scale language modeling, we apply a similar approach towards building a single generalist agent beyond the realm of text outputs. The agent, which we refer to as Gato, works as a multi-modal, multi-task, multi-embodiment generalist policy. The same network with the same weights can play Atari, caption images, chat, stack blocks with a real robot arm and much more, deciding based on its context whether to output text, joint torques, button presses, or other tokens. In this report we describe the model and the data, and document the current capabilities of Gato.

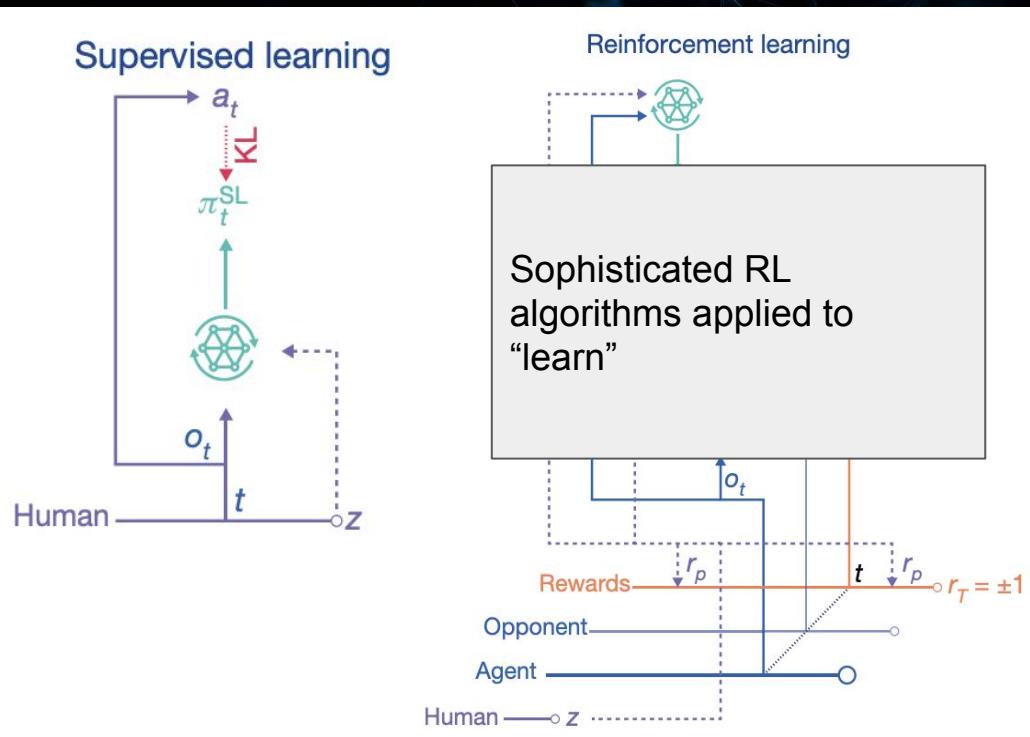


AlphaStar 2019



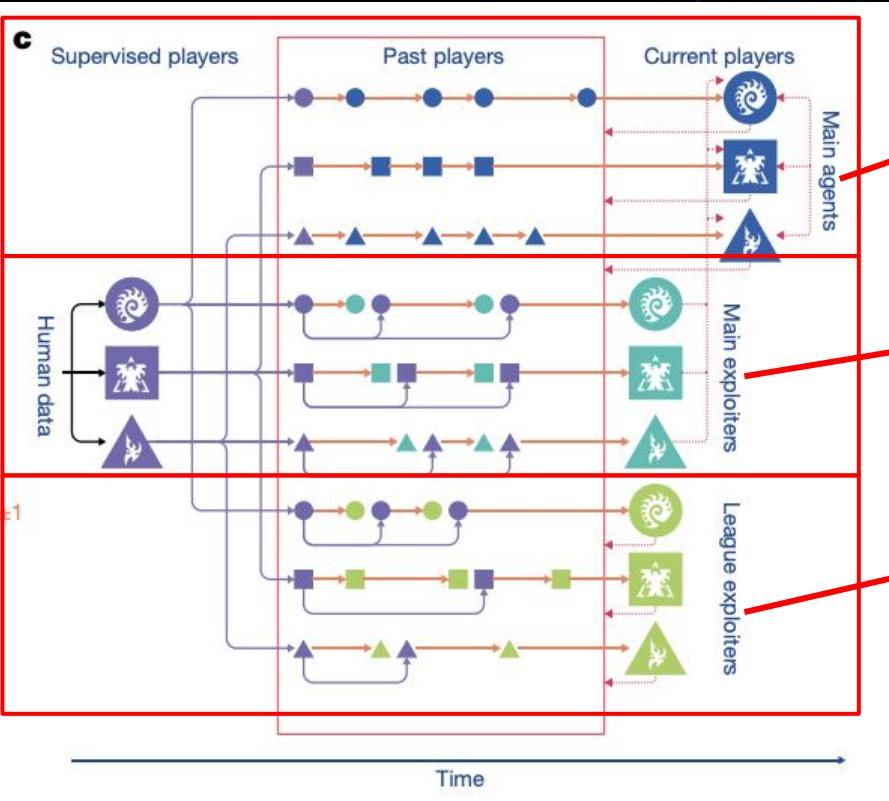
- Starcraft is more complex than board games like Go
 - Go has $\sim 19^*19$ action spaces
 - Starcraft has huge (thus sparse) action spaces!
 - E.g. image pixel level
 - Or “what, who, where, when” action sequence!
- To make things more challenging!
 - Imperfect info (e.g. no direct visibility to opponent)
 - AI has to respond in ~100ms

AlphaStar 2019, combine supervised and reinforcement learning (remember alphago?)



- Imitate experts by learning human replays, conditioned on z (i.e. build order stats)
- Off-policy on learning from “experience replay” using a2c and other algorithms
- “exploration of novel strategy” vs exploitation vs “robustness”
 - more reward toward supervised policy

AlphaStar 2019, multi-agent league training



3 groups, 12 agents in total

- Main agents (1 x 3)

- Self-play with all other 3 groups of agents and try to improve

- Main exploiters (1 x 3)

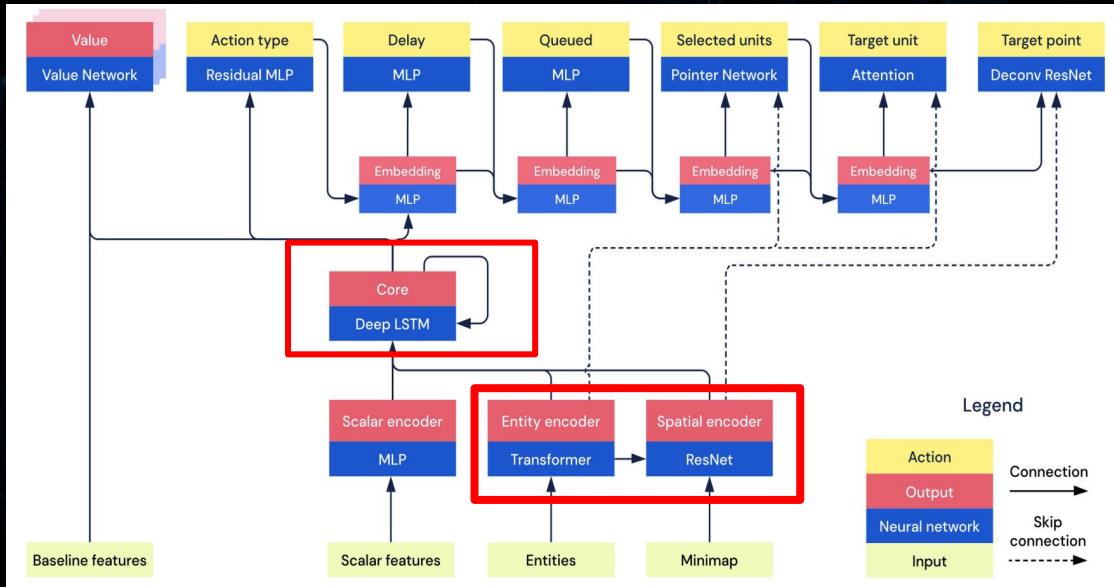
- Exploit weakness in “Main” group by playing with Main agents only

- League exploiters (2 x 3)

- Try to find systematic weakness of the “entire league”

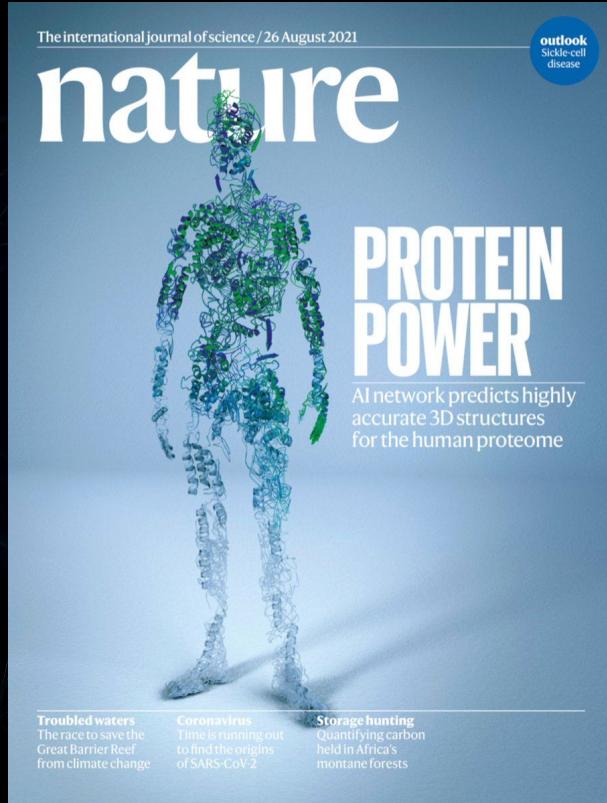
Exploiters periodically re-initialized (purple shapes) to increase diversity and robustness

AlphaStar 2019, network architecture



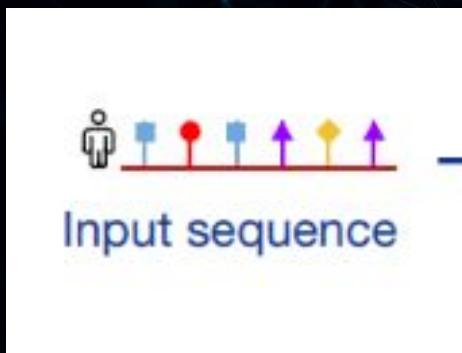
- Notice the network output is structured
 - What: Action type
 - When: Delay/Queued
 - Who: Selected unit and Target units
 - Where: Target point
- LSTM (RNN) used to process sequence
- ResNet for minimap
- Transformer to apply “self attention” on units/entities

AlphaFold 2021

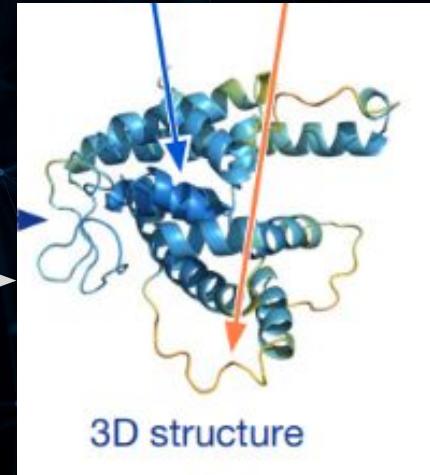


- Protein folding 3d structure prediction is a 50-year scientific problem!
 - Without AlphaFold, people rely on expensive and tedious work to observe and record every protein shape
 - AlphaFold completely changed the game to get atomic level precision!
- Transformer alike attention model applied

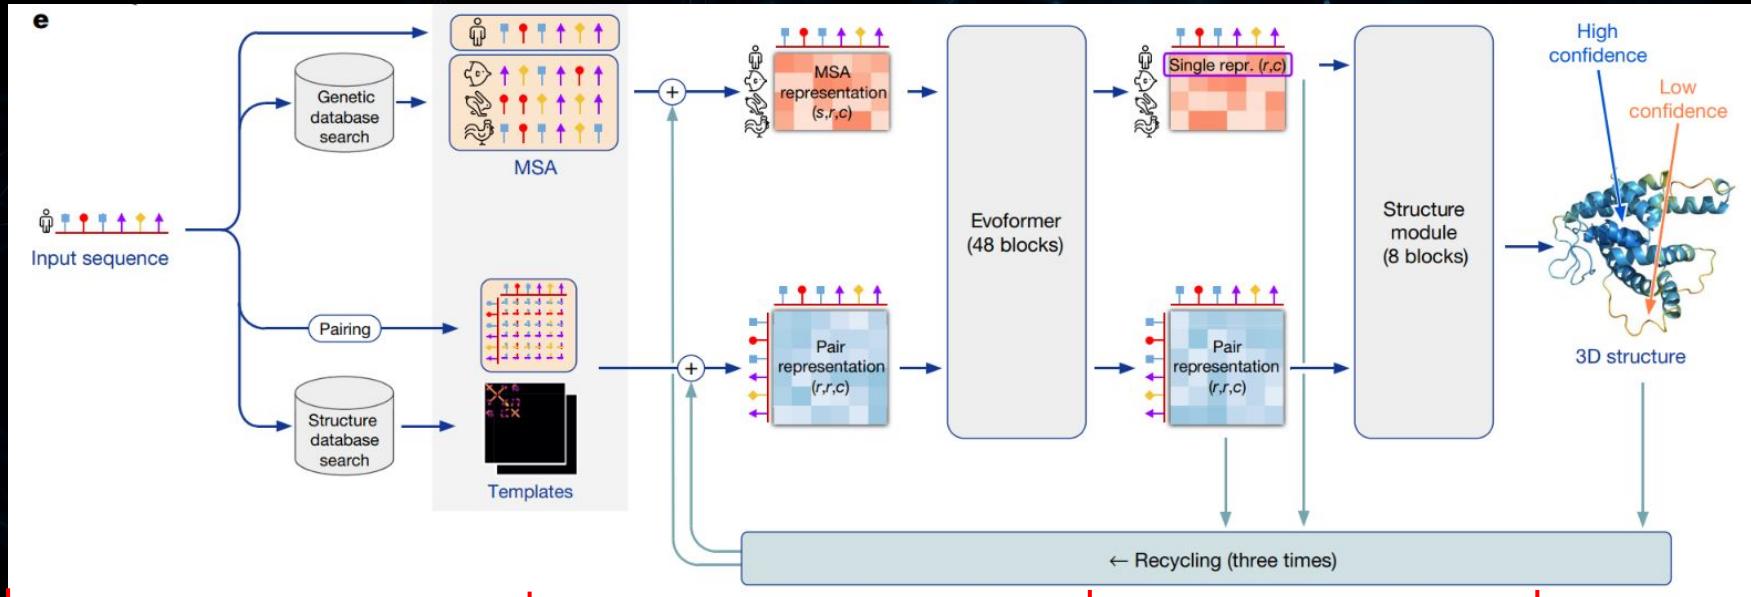
AlphaFold 2021, problem in short



→ How?



AlphaFold 2021, main model structure



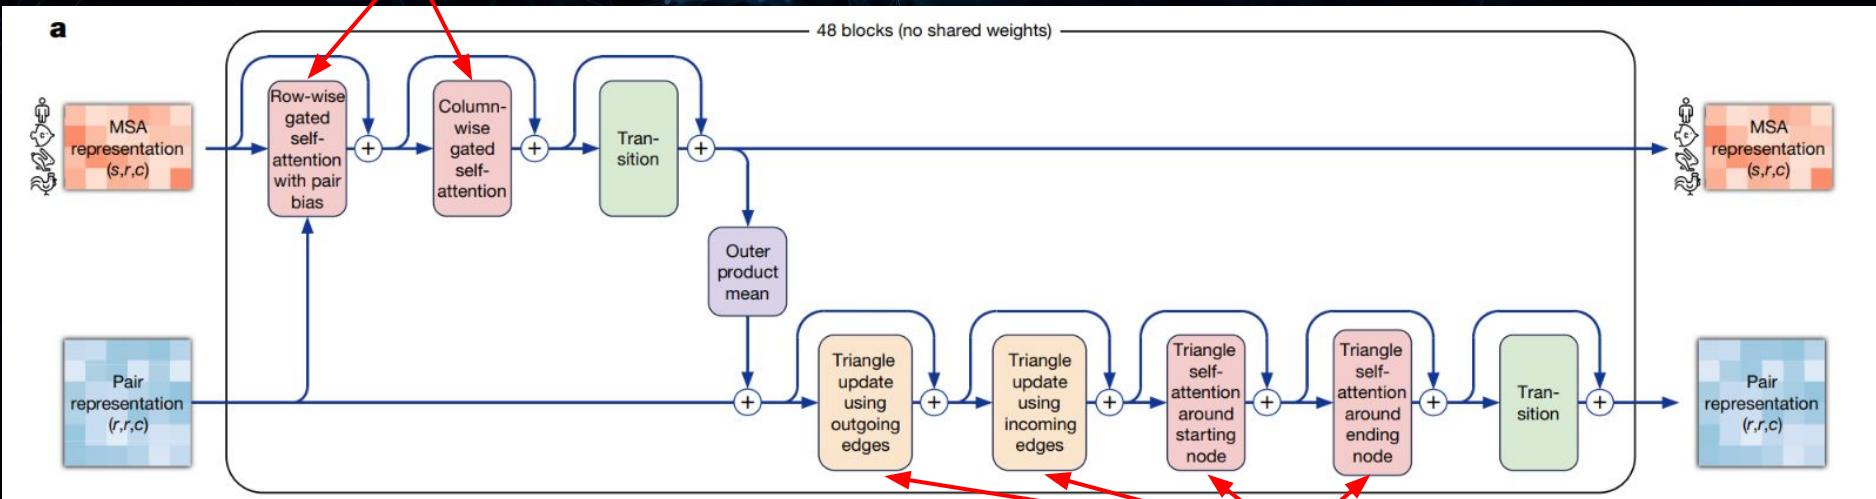
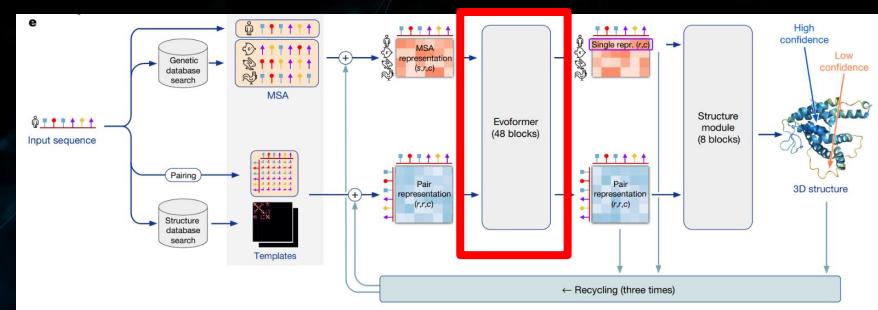
Feature Engineering

Feature fusion and
Encoding

Decoding

AlphaFold 2021, Evoformer

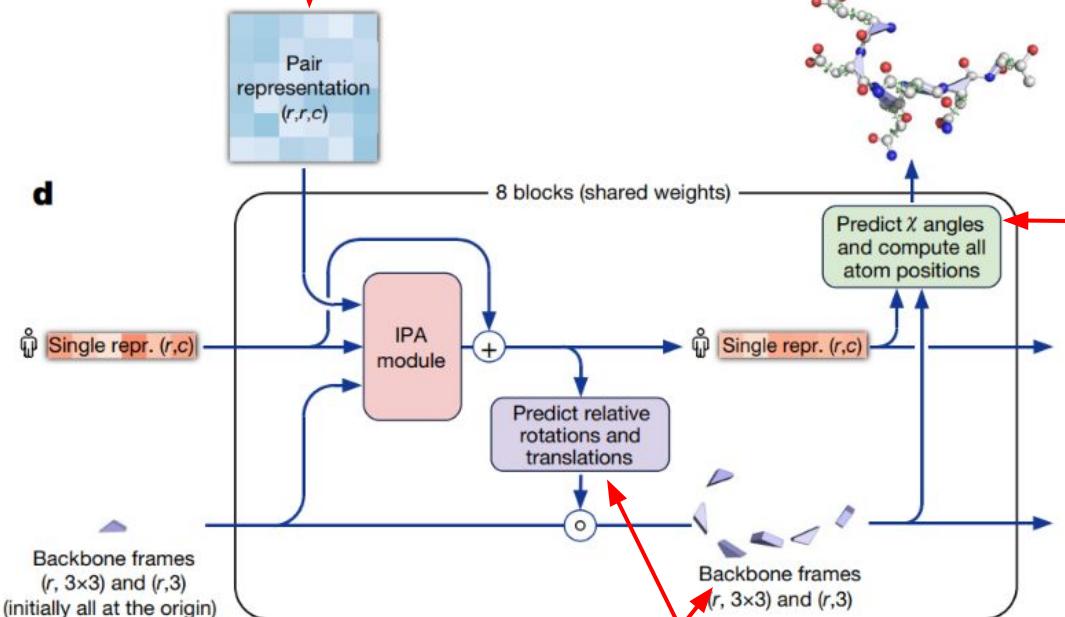
Row and column wise gated attention for 2D data



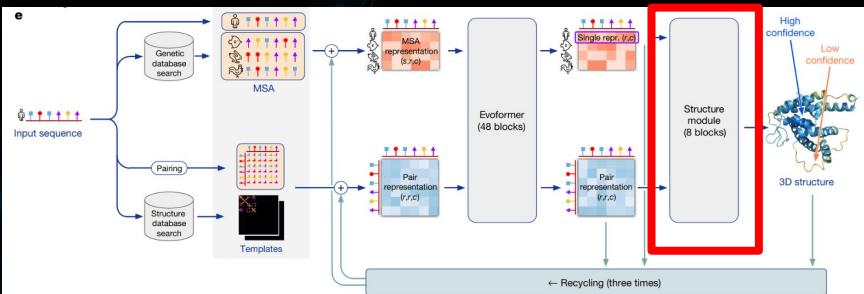
- Triangle Update? Think about triangle as stable structure in 3d space, and $d_1+d_2 > d_3$
- Directed graph, thus treat outgoing/incoming differently

AlphaFold 2021, Structure Module

Pair rep from Evoformer output

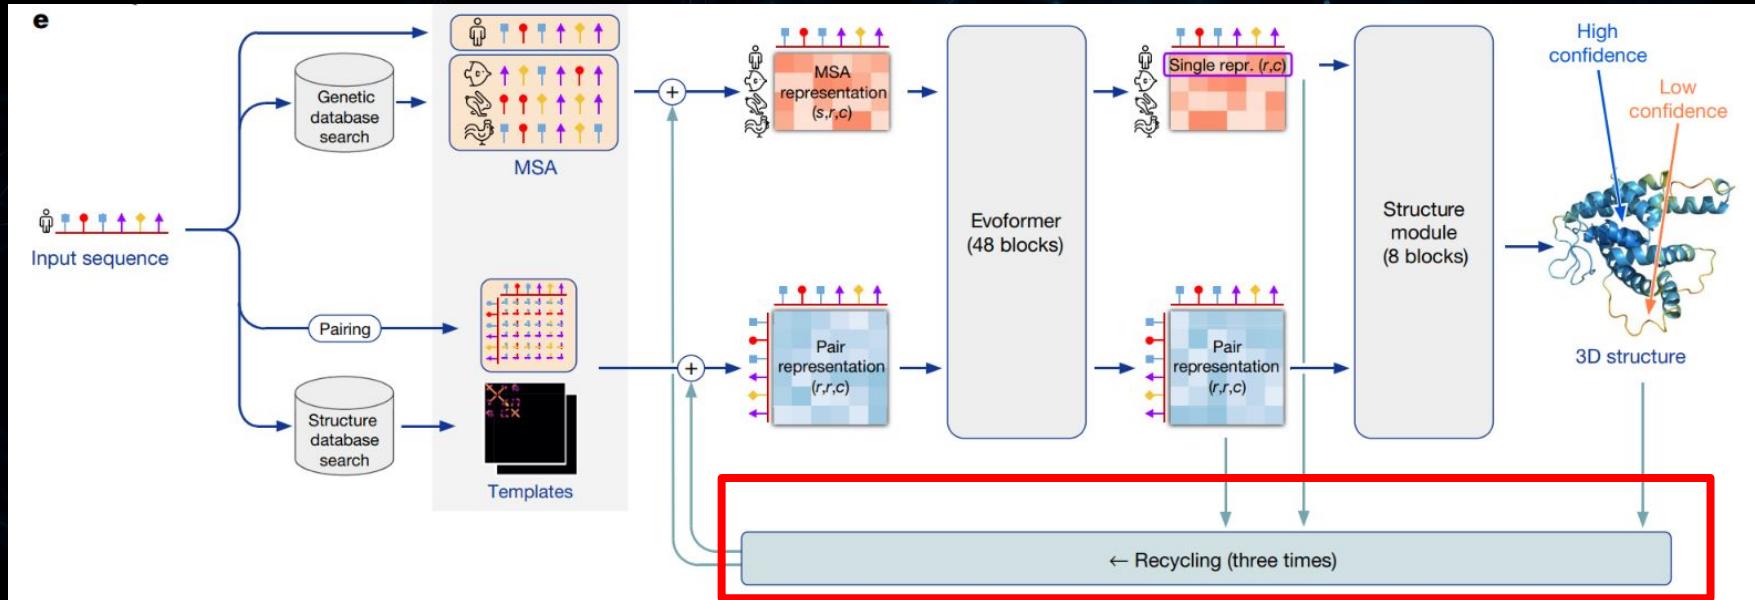


(3x3) vector as rotation, and (3) vector as position translation



With rotation and translation
vectors, build 3D representation

AlphaFold 2021, 3x “recycling” (similar to deeper NN)



Equivalent to a 3x deep network
with residual connection

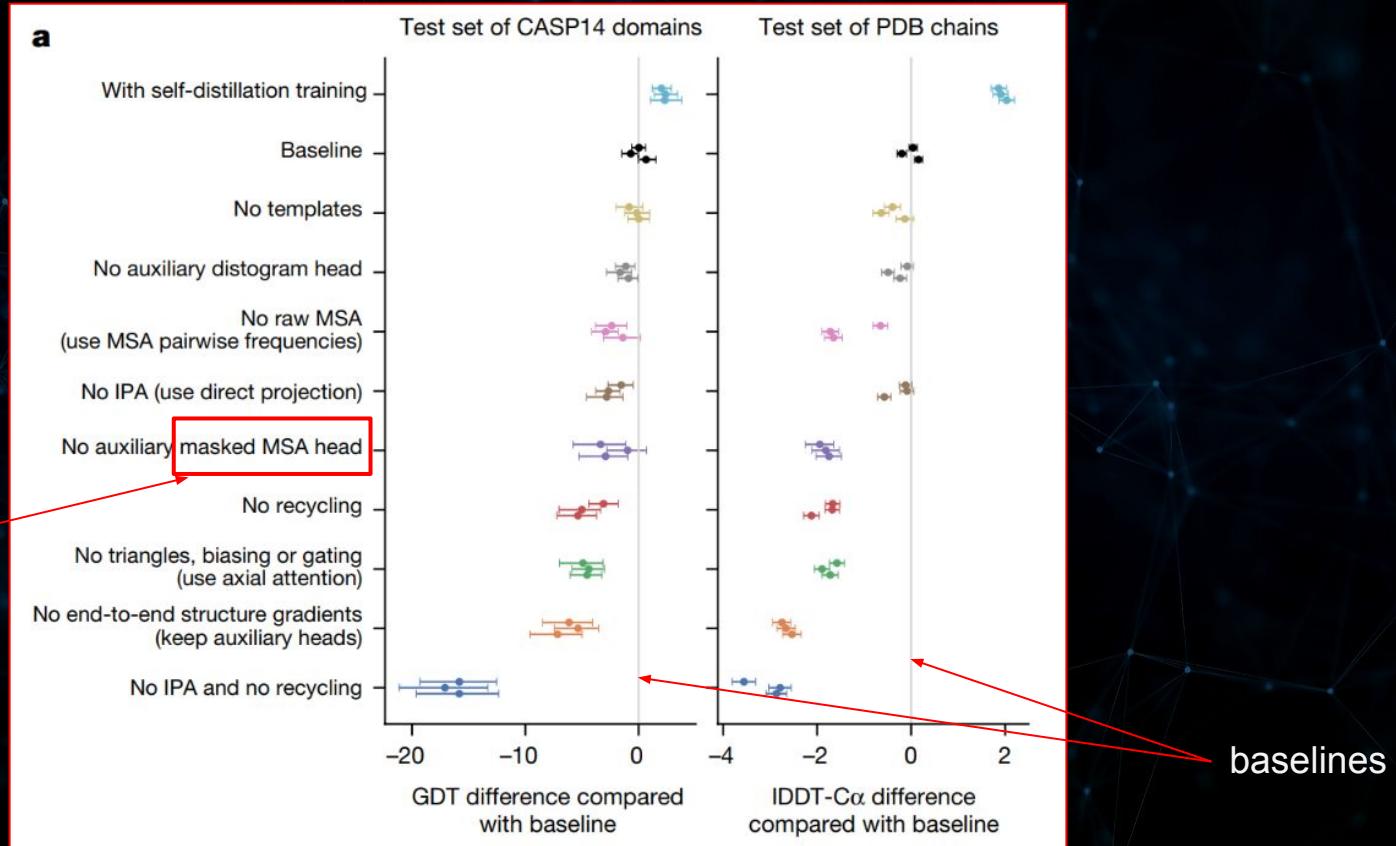
AlphaFold 2021, other training tricks

Training with labelled and unlabelled data

The AlphaFold architecture is able to train to high accuracy using only supervised learning on PDB data, but we are able to enhance accuracy (Fig. 4a) using an approach similar to noisy student self-distillation³⁵. In this procedure, we use a trained network to predict the structure of around 350,000 diverse sequences from Uniclust30³⁶ and make a new dataset of predicted structures filtered to a high-confidence subset. We then train the same architecture again from scratch using a mixture of PDB data and this new dataset of predicted structures as the training data, in which the various training data augmentations such as cropping and MSA subsampling make it challenging for the network to recapitulate the previously predicted structures. This self-distillation procedure makes effective use of the unlabelled sequence data and considerably improves the accuracy of the resulting network.

AlphaFold 2021, model/training feature ablation

BERT-alike
pretraining is
everywhere



AlphaCode 2022



2022-3-16

Competition-Level Code Generation with AlphaCode

Yujia Li^{*}, David Choi^{*}, Junyoung Chung^{*}, Nate Kushman^{*}, Julian Schrittwieser^{*}, Rémi Leblond^{*}, Tom Eccles^{*}, James Keeling^{*}, Félix Gimeno^{*}, Agustin Dal Lago^{*}, Thomas Hubert^{*}, Peter Choy^{*}, Cyprien de Masson d'Autume^{*}, Igor Babuschkin, Xinyun Chen, Po-Sen Huang, Johannes Welbl, Sven Gowal, Alexey Cherepanov, James Molloy, Daniel J. Mankowitz, Esme Sutherland Robson, Pushmeet Kohli, Nando de Freitas, Koray Kavukcuoglu and Oriol Vinyals
*Joint first authors

Programming is a powerful and ubiquitous problem-solving tool. Developing systems that can assist programmers or even generate programs independently could make programming more productive and accessible, yet so far incorporating innovations in AI has proven challenging. Recent large-scale language models have demonstrated an impressive ability to generate code, and are now able to complete simple programming tasks. However, these models still perform poorly when evaluated on more complex, unseen problems that require problem-solving skills beyond simply translating instructions into code. For example, competitive programming problems which require an understanding of algorithms and complex natural language remain extremely challenging. To address this gap, we introduce AlphaCode, a system for code generation that can create novel solutions to these problems that require deeper reasoning. In simulated evaluations on recent programming competitions on the Codeforces platform, AlphaCode achieved on average a ranking of top 54.3% in competitions with more than 5,000 participants. We found that three key components were critical to achieve good and reliable performance: (1) an extensive and clean competitive programming dataset for training and evaluation, (2) large and efficient-to-sample transformer-based architectures, and (3) large-scale model sampling to explore the search space, followed by filtering based on program behavior to a small set of submissions.

1	Introduction	2	5.4 Results on APPS	19
2	Problem setup	5	6 AlphaCode's capabilities & limitations	20
2.1	Competitive programming	5	6.1 Copying from training data	21
2.2	Evaluation	6	6.2 Model solution characteristics	22
3	Datasets	6	6.3 Sensitivity to problem descriptions	24
3.1	Pre-training dataset	7	6.4 Sensitivity to provided metadata	24
3.2	CodeContests fine-tuning dataset	7	6.5 Loss is a poor proxy for solve rate	26
4	Approach	9	7 Related work	27
4.1	Model architecture	10	7.1 Program synthesis	27
4.2	Pre-training	11	7.2 Transformers for program synthesis	28
			7.3 Scaling sampling	28

- One of the powerful coding LLMs
 - GPT3 (codex), T5, LaMDA, PaLM...
- Focus on the very hard “*Competitive coding*” problem
 - GPT3 or codex alike AI are more “instruction to code” translation
 - AlphaCode targets to understand a long and difficult problem statement and provide a solution
- Achieved 54.3% competitor level

AlphaCode 2022, architecture and training

- Encoder-decoder transformer
 - Natural language seq, to expect code output seq
 - Encoder output can be cached for sampling later
- Pre-training
 - “Masked language modelling loss”, similar to [BERT](#), on Github dataset
 - [AdamW](#) with weight decay and gradient clipping
 - *Scaling: [2018 OpenAI scaling rule](#) is used, but [Chinchilla paper](#) recommends different scaling settings*
- Fine-tuning
 - CodeContext dataset (the format of the problems to be solved)
 - Focus on precision because one problem has multiple solutions!
 - Ignore missed token penalty if not in its distribution
 - Tags added to input text (prompt), e.g. “dp” (dynamic programming), “bfs” (breadth-first search), and rating (e.g. 1200 points)

AlphaCode 2022, sampling, filtering & clustering

- Sampling
 - Millions of samples for each problem! Half in python, half in c++
 - “Randomize problem tags” (e.g. dp, bfs, dfs) “and ratings” (e.g. 1200)
 - As if a human competitor guesses which difficulty level the problem is and what high-level approach to take
 - Set high temperature (less confident, but encourage diversity) and apply sampling decoding
- Filtering
 - Run the generated code and keep those that passed sample tests
 - Still there can be 10k+ solutions after filtering
 - Or, 10% problems have no single solution passed with 1M+ sampling
- Clustering
 - Additional test inputs added, group the samples by additional outputs
 - Thus a separate test-input-generation model is needed!
 - Select one solution from largest to smallest clusters afterward

AlphaTensor 2022



- The core problem
 - Multiplication is more expensive than addition in GPU/TPU/CPU
 - Fewer multiplications (but more additions/subtractions) to achieve faster computation
 - AI to discover algorithms to reduce multiplications for matrix multiplication
 - Design a game to train this AI

Let's review the intuition and details!

AlphaTensor 2022 (Intuition, borrowed from [Yannic Youtube video](#))

- Let's use a related simple 1-dimension example

$$a^2 - b^2$$

- 2 multiplications and 1 addition/subtraction
 - multiplication: $h_1 = a \times a$
 - multiplication: $h_2 = b \times b$
 - addition/subtraction: $h_1 - h_2$
 - But we have a way to reduce the number of multiplications!
 - We change rewrite.
 $a^2 - b^2 = (a+b) \times (a-b)$
 - 2 addition/subtraction, and 1 multiplication!
 - Can we apply similar idea to matrix multiplication?
 - AlphaTensor: Yes!

AlphaTensor 2022 (2x2 matrix mul, mapping to 4x4x4)

$$\begin{pmatrix} c_1 & c_2 \\ c_3 & c_4 \end{pmatrix} = \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \cdot \begin{pmatrix} b_1 & b_2 \\ b_3 & b_4 \end{pmatrix}$$

$$c_1 = a_1 \times b_1 + a_2 \times b_3$$

$$c_2 = a_1 \times b_2 + a_2 \times b_4$$

$$c_3 = a_3 \times b_1 + a_4 \times b_3$$

$$c_4 = a_3 \times b_2 + a_4 \times b_4$$

- 8 multiplications!

AlphaTensor 2022 (2x2 matrix mul, mapping to 4x4x4)

$$\begin{pmatrix} c_1 & c_2 \\ c_3 & c_4 \end{pmatrix} = \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \cdot \begin{pmatrix} b_1 & b_2 \\ b_3 & b_4 \end{pmatrix}$$

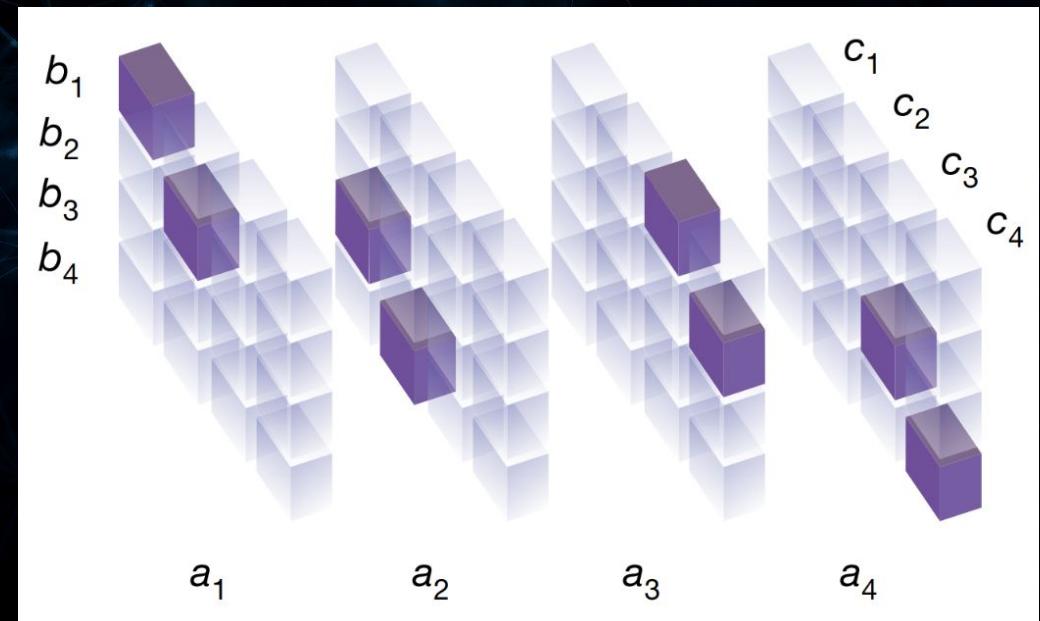
$$c_1 = a_1 \times b_1 + a_2 \times b_3$$

$$c_2 = a_1 \times b_2 + a_2 \times b_4$$

$$c_3 = a_3 \times b_1 + a_4 \times b_3$$

$$c_4 = a_3 \times b_2 + a_4 \times b_4$$

- 8 multiplications!



AlphaTensor 2022 (2x2 matrix mul, mapping to 4x4x4)

$$\begin{pmatrix} c_1 & c_2 \\ c_3 & c_4 \end{pmatrix} = \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \cdot \begin{pmatrix} b_1 & b_2 \\ b_3 & b_4 \end{pmatrix}$$

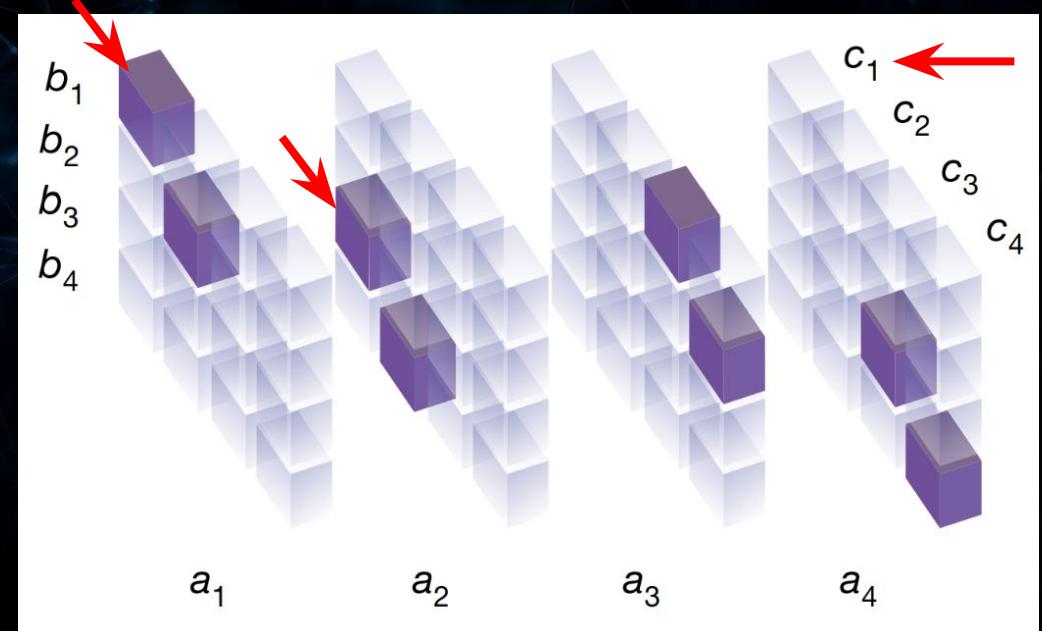
$$c_1 = a_1 \times b_1 + a_2 \times b_3$$

$$c_2 = a_1 \times b_2 + a_2 \times b_4$$

$$c_3 = a_3 \times b_1 + a_4 \times b_3$$

$$c_4 = a_3 \times b_2 + a_4 \times b_4$$

- 8 multiplications!



AlphaTensor 2022 (2x2 matrix mul, mapping to 4x4x4)

$$\begin{pmatrix} c_1 & c_2 \\ c_3 & c_4 \end{pmatrix} = \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \cdot \begin{pmatrix} b_1 & b_2 \\ b_3 & b_4 \end{pmatrix}$$

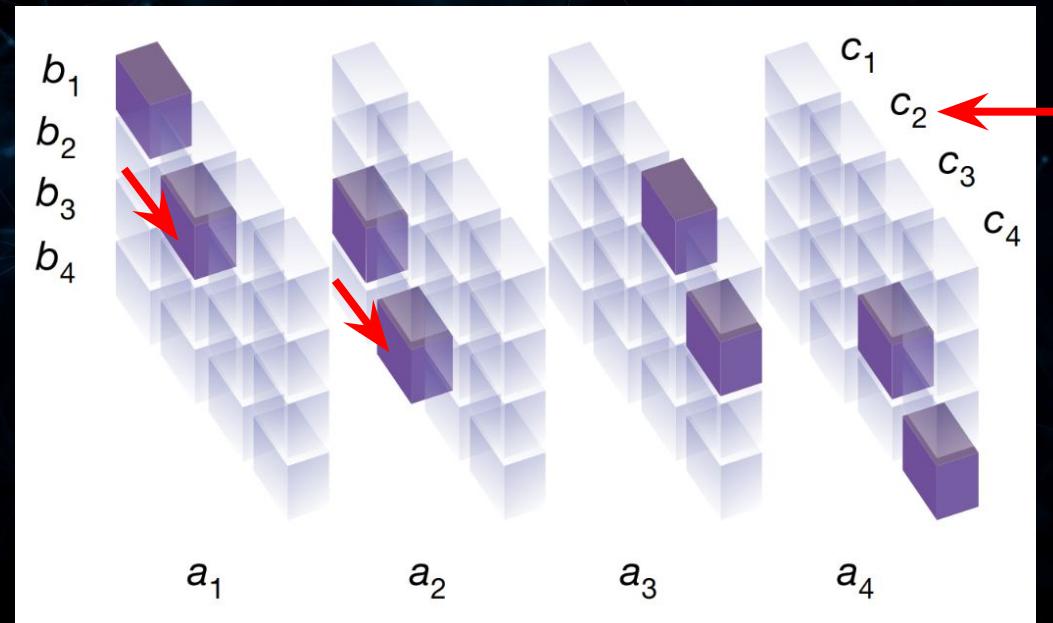
$$c_1 = a_1 \times b_1 + a_2 \times b_3$$

$$c_2 = a_1 \times b_2 + a_2 \times b_4$$

$$c_3 = a_3 \times b_1 + a_4 \times b_3$$

$$c_4 = a_3 \times b_2 + a_4 \times b_4$$

- 8 multiplications!



AlphaTensor 2022 (Strassen's algorithm, 7 multiplications!)

$$\begin{pmatrix} c_1 & c_2 \\ c_3 & c_4 \end{pmatrix} = \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \cdot \begin{pmatrix} b_1 & b_2 \\ b_3 & b_4 \end{pmatrix}$$

$$c_1 = a_1 \times b_1 + a_2 \times b_3$$

$$c_2 = a_1 \times b_2 + a_2 \times b_4$$

$$c_3 = a_3 \times b_1 + a_4 \times b_3$$

$$c_4 = a_3 \times b_2 + a_4 \times b_4$$

- 8 multiplications!

$$m_1 = (a_1 + a_4)(b_1 + b_4)$$

$$m_2 = (a_3 + a_4)b_1$$

$$m_3 = a_1(b_2 - b_4)$$

$$m_4 = a_4(b_3 - b_1)$$

$$m_5 = (a_1 + a_2)b_4$$

$$m_6 = (a_3 - a_1)(b_1 + b_2)$$

$$m_7 = (a_2 - a_4)(b_3 + b_4)$$

$$c_1 = m_1 + m_4 - m_5 + m_7$$

$$c_2 = m_3 + m_5$$

$$c_3 = m_2 + m_4$$

$$c_4 = m_1 - m_2 + m_3 + m_6$$

AlphaTensor 2022 (gamification)

$$m_1 = (a_1 + a_4)(b_1 + b_4)$$

$$m_2 = (a_3 + a_4)b_1$$

$$m_3 = a_1(b_2 - b_4)$$

$$m_4 = a_4(b_3 - b_1)$$

$$m_5 = (a_1 + a_2)b_4$$

$$m_6 = (a_3 - a_1)(b_1 + b_2)$$

$$m_7 = (a_2 - a_4)(b_3 + b_4)$$

$$c_1 = m_1 + m_4 - m_5 + m_7$$

$$c_2 = m_3 + m_5$$

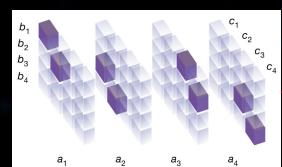
$$c_3 = m_2 + m_4$$

$$c_4 = m_1 - m_2 + m_3 + m_6$$

- u represents weighted sum of a_i
- v represents weighted sum of b_j
- w represents weighted sum of intermediate multiplications (landscape view)
- **outer_product** (u, v, w) forms a $4 \times 4 \times 4$ matrix for each rank and requires 1 multiplication and multiple additions

Game:

Find $\text{rank} \leq n^3$ $u/v/w$ combinations, so that sum of sub-matrix is equal to expected



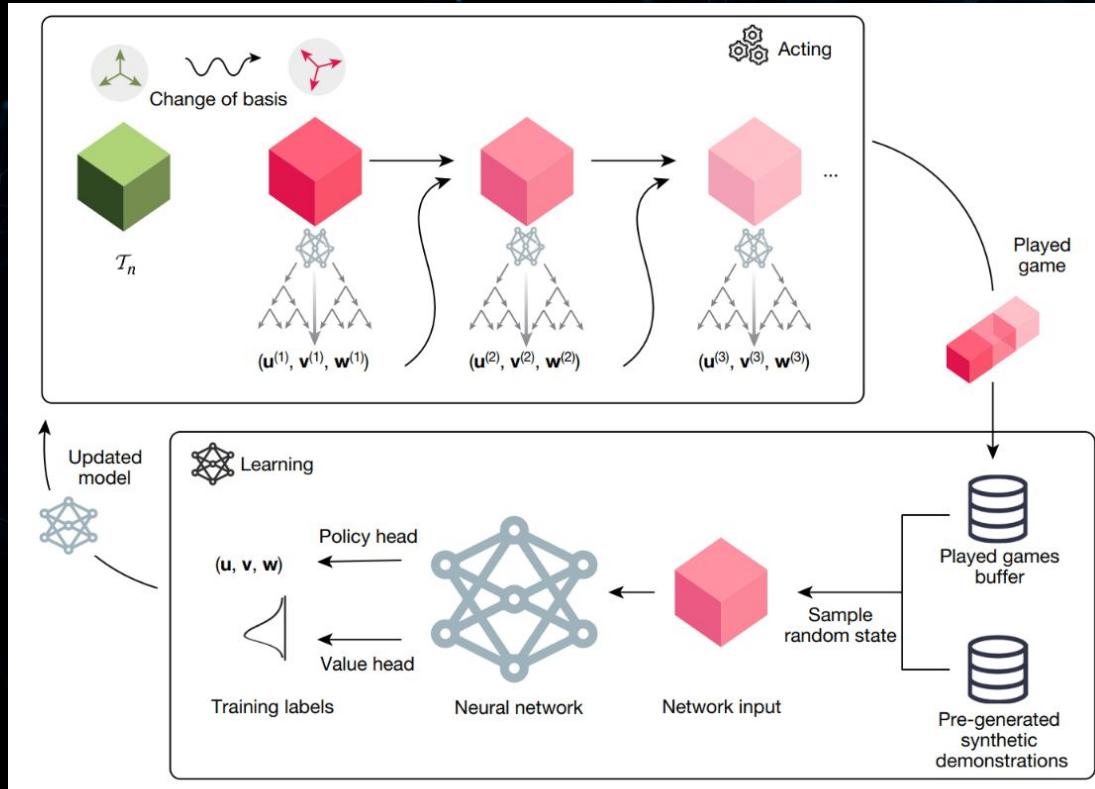
$$\mathbf{U} = \begin{pmatrix} 1 & 0 & 1 & 0 & 1 & -1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & -1 \end{pmatrix}$$

$$\mathbf{V} = \begin{pmatrix} 1 & 1 & 0 & -1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & -1 & 0 & 1 & 0 & 1 \end{pmatrix}$$

$$\mathbf{W} = \begin{pmatrix} 1 & 0 & 0 & 1 & -1 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & -1 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

$$\mathcal{T}_n = \sum_{r=1}^R \mathbf{u}^{(r)} \otimes \mathbf{v}^{(r)} \otimes \mathbf{w}^{(r)}$$

AlphaTensor 2022 (model arch)



Similar to AlphaZero

Self-play to try to win the game with lower rank (fewer multiplications)

Data augmentation through data synthesis, change of matrix basis, and rank-order shuffling

AlphaTensor 2022 (oversimplified impact description) borrowed from [Dr. Hung-yi Lee video](#)

Block Multiplication

$$A = \begin{matrix} \text{Gray} & \text{Blue} \\ \text{Orange} & \text{Gray} \end{matrix} \quad n \times n \quad n/2 \times n/2$$

"2 x 2"

$$B = \begin{matrix} \text{Yellow} & \text{Blue} \\ \text{Green} & \text{Gray} \end{matrix} \quad n \times n$$

$$AB = \left[\begin{array}{c} \begin{matrix} \text{Gray} & \text{Yellow} \\ \text{Orange} & \text{Yellow} \end{matrix} + \begin{matrix} \text{Blue} & \text{Green} \\ \text{Gray} & \text{Green} \end{matrix} & \begin{matrix} \text{Gray} & \text{Blue} \\ \text{Orange} & \text{Blue} \end{matrix} + \boxed{\begin{matrix} \text{Blue} & \text{Gray} \\ \text{Blue} & \text{Gray} \end{matrix}} \\ \hline \begin{matrix} \text{Orange} & \text{Yellow} \\ \text{Gray} & \text{Green} \end{matrix} + \begin{matrix} \text{Gray} & \text{Green} \\ \text{Gray} & \text{Green} \end{matrix} & \begin{matrix} \text{Orange} & \text{Blue} \\ \text{Gray} & \text{Gray} \end{matrix} + \begin{matrix} \text{Gray} & \text{Gray} \\ \text{Gray} & \text{Gray} \end{matrix} \end{array} \right]$$

8 matrix multiplications? No, you only need **7**.

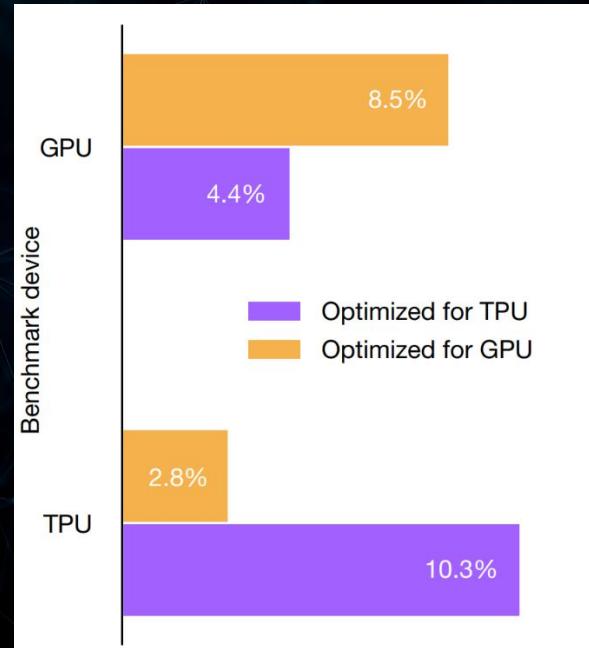
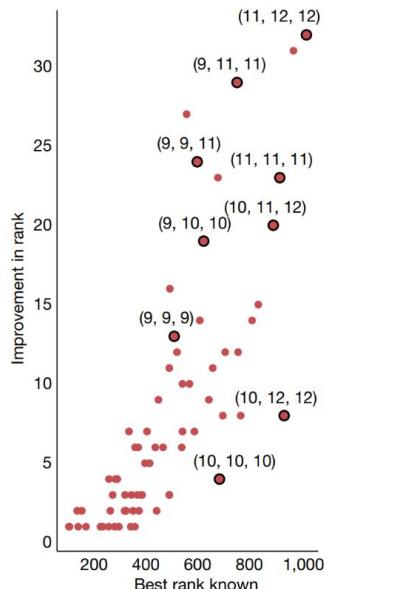
Strassen algorithm

You only need 7
multiplications as well.

Recursively

AlphaTensor 2022 (impact)

Size (n, m, p)	Best method known	Best rank known	AlphaTensor rank Modular Standard
(2, 2, 2)	(Strassen, 1969) ²	7	7
(3, 3, 3)	(Laderman, 1976) ¹⁵	23	23
(4, 4, 4)	(Strassen, 1969) ²	49	47
(5, 5, 5)	(3, 5, 5) + (2, 5, 5)	98	96
(2, 2, 3)	(2, 2, 2) + (2, 2, 1)	11	11
(2, 2, 4)	(2, 2, 2) + (2, 2, 2)	14	14
(2, 2, 5)	(2, 2, 2) + (2, 2, 3)	18	18
(2, 3, 3)	(Hopcroft and Kerr, 1971) ¹⁶	15	15
(2, 3, 4)	(Hopcroft and Kerr, 1971) ¹⁶	20	20
(2, 3, 5)	(Hopcroft and Kerr, 1971) ¹⁶	25	25
(2, 4, 4)	(Hopcroft and Kerr, 1971) ¹⁶	26	26
(2, 4, 5)	(Hopcroft and Kerr, 1971) ¹⁶	33	33
(2, 5, 5)	(Hopcroft and Kerr, 1971) ¹⁶	40	40
(3, 3, 4)	(Smirnov, 2013) ¹⁸	29	29
(3, 3, 5)	(Smirnov, 2013) ¹⁸	36	36
(3, 4, 4)	(Smirnov, 2013) ¹⁸	38	38
(3, 4, 5)	(Smirnov, 2013) ¹⁸	48	47
(3, 5, 5)	(Sedoglavic and Smirnov, 2021) ¹⁹	58	58
(4, 4, 5)	(4, 4, 2) + (4, 4, 3)	64	63
(4, 5, 5)	(2, 5, 5) \otimes (2, 1, 1)	80	76



Algorithm Search!

Device-specific optimization

- A different reward function to include device benchmark used

Some of my personal thoughts

- Breakthrough research from Deepmind is built on top of
 - Vision and support from leads
 - Top research and excellent engineering
 - Scalable ML infra by Google
- A shift from games to science/engineering?
 - Alpha[*] work focused on fun games before 2019
 - Biology, engineering, math/algorithm since 2020
- The return of Reinforcement Learning 2022

References

Papers

- [AlphaGo](#) (nature)
- [AlphaGo Zero](#) (nature)
- [AlphaZero](#) (science)
- [AlphaStar](#) (science)
- [AlphaFold2](#) (nature)
- [AlphaCode](#) (arxiv)
- [AlphaTensor](#) (nature)

Other references

- [Jonathan Hui: AlphaGo: How it works technically?](#)
- [Yannic Kilcher - YouTube](#)
- [Mu Li - YouTube](#)
- [Shusen Wang - YouTube](#)
- [Two Minute Papers - YouTube](#)
- [Hung-yi Lee - YouTube](#)

Thank you!

fun-ai-talks

Let's read 7 DeepMind Alpha* papers!

AlphaGo

2016

AlphaGo Zero

2017

AlphaZero

2018

AlphaStar

2019

AlphaFold 2

2021

AlphaCode

2022

AlphaTensor

hululu.zhu@gmail.com

10/2022